

# How much is your car worth? A Used Car Price Prediction System (UCPPS)

Faculty of Graduate Studies & Research

Instructor: Dr. Alireza Manashty

Student Name: Tanu Nanda Prabhu

4/8/2020



University  
of Regina

# Table of Contents

- Introduction
- Problem Statement
- Data
- Model Planning and building
- Solution Overview
- Outcome
- Tools
- Team Roles
- Time line
- References

# INTRODUCTION

- Vehicle value forecast is a significant errand particularly when the vehicle is used.
- The value of the car depends on several factors:
  - Make (brand of the car)
  - Power
  - Number of kilometers it has been run
  - Year of registration, and many more
- Better the features higher the price



Technical Data		
<b>Cylinders / Capacity (cc)</b> In-line 4 / 1,998	<b>Petrol / Diesel</b> Petrol	<b>Transmission type</b> 8-speed Steptronic Sport transmission
<b>Combustion Engine</b> Max output (kW/hp/rpm) 185 / 252 / 5,200 - 6,500	<b>Max torque (Nm/rpm)</b> 350 / 1,450 - 4,800	<b>Acceleration 0 - 100km/h (s)</b> 6.2
<b>Top speed (km/h)</b> 250	<b>Fuel consumption (ltr/100km)</b> 5.8	<b>CO<sub>2</sub> emissions (g/km)</b> 132
Manufacturer Recommended Net Selling Price		RM 398,071.00
Personal Registration		RM 350.00
Registration Fees & HP Endorsement		RM 379.00
Road Tax		RM 398,800.00
Recommended Retail Price without Insurance**		

**5 YEARS**  
**BMW WARRANTY**  
WITH FREE SCHEDULED SERVICE



\* Actual car specifications may vary from the picture shown above.  
\* Prices reflect those as of 15 April 2017. Based on Manufacturer Recommended Net Selling Price, and only valid in Peninsular Malaysia.  
\* Prices include 6% SST.  
\* Prices and specifications are subject to change without prior notice.  
\* BMW Malaysia is a member of BMW Group Malaysia.  
\*\* Financing rates are based on an estimated 3.99% interest rate. For more information about BMW Financing Solutions, please contact BMW Credit at 1800 88 3000. Terms and conditions apply.  
\*\*\* All retail prices are inclusive of BMW 5 Year Unlimited Mileage Warranty + Free Service.  
\*\*\* All retail prices are inclusive of BMW 5 Year Warranty, valid for 100,000 km or more, whichever is longer.

BMW Voice: 1800 88 3000  
www.bmw.com.my


Image Credits: [Manashty, 2020](#) [2]

# PROBLEM STATEMENT

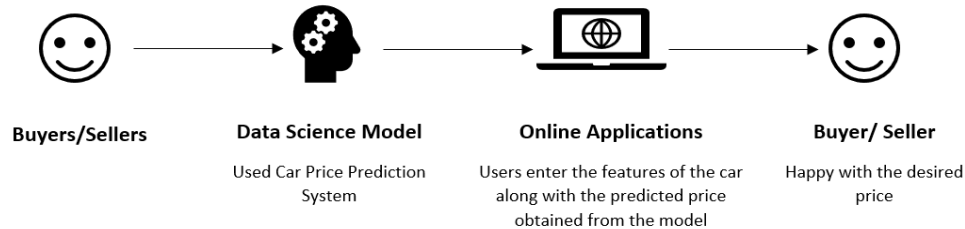
- Used Car Prices are important reflection of the economy and they greatly interest both buyers and sellers.
- A prediction model that estimates resale price based on car's attributes or features is much more needed today.
- My analysis aims to determine which features of the car that may have the strongest statistical correlation with the price of the car.

# Problem Statement

## •Current Situation



## •Desired Situation



# DATA

- The data set was chosen from [data.world](https://data.world), data was originally scraped from e-bay [3]

```
import pandas as pd
import time
start_time = time.time()
df = pd.read_csv("/content/drive/My Drive/Dataset/autos.csv", sep = ',', header = 0, encoding='cp1252')
print("--- %s seconds ---" % (time.time() - start_time))
df.head(5)
```

--- 9.454026222229004 seconds ---

	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model
0	2016-03-24 11:52:17	Golf_3_1.6	privat	Angebot	480	test	NaN	1993	manuell	0	golf
1	2016-03-24 10:58:45	A5_Sportback_2.7_Tdi	privat	Angebot	18300	test	coupe	2011	manuell	190	NaN
2	2016-03-14 12:52:21	Jeep_Grand_Cherokee_"Overland"	privat	Angebot	9800	test	suv	2004	automatik	163	grand
3	2016-03-17 16:54:04	GOLF_4_1_4__3TÜRER	privat	Angebot	1500	test	kleinwagen	2001	manuell	75	golf
4	2016-03-31 17:25:20	Skoda_Fabia_1.4_TDI_PD_Classic	privat	Angebot	3600	test	kleinwagen	2008	manuell	69	fabia

# DATA STATISTICS

```
df.info() # Getting information about the datatypes
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 371528 entries, 0 to 371527
Data columns (total 20 columns):
#   Column              Non-Null Count  Dtype
---  -
0   dateCrawled          371528 non-null object
1   name                 371528 non-null object
2   seller               371528 non-null object
3   offerType            371528 non-null object
4   price                371528 non-null int64
5   abtest               371528 non-null object
6   vehicleType          333659 non-null object
7   yearOfRegistration   371528 non-null int64
8   gearbox              351319 non-null object
9   powerPS              371528 non-null int64
10  model                351044 non-null object
11  kilometer            371528 non-null int64
12  monthOfRegistration   371528 non-null int64
13  fuelType             338142 non-null object
14  brand                371528 non-null object
15  notRepairedDamage    299468 non-null object
16  dateCreated           371528 non-null object
17  nrOfPictures          371528 non-null int64
18  postalCode            371528 non-null int64
19  lastSeen              371528 non-null object
dtypes: int64(7), object(13)
memory usage: 56.7+ MB
```

- **dateCrawled** : when this ad was first crawled, all field-values are taken from this date
- **name** : 'name' of the car
- **seller** : private or dealer
- **offerType** : With offer or without offer
- **price** : the price on the ad to sell the car
- **abtest** : Test on the car
- **vehicleType** : Type of the car (Sedan, truck, etc.)
- **yearOfRegistration** : at which year the car was first registered
- **gearbox** : Automatic or manual transmission
- **powerPS** : power of the car in PS
- **model** : Model of the car
- **kilometer** : how many kilometers the car has driven
- **monthOfRegistration** : at which month the car was first registered
- **fuelType** : Gas, Petrol, Diesel, etc.
- **brand** : Mercedes, Audi, BMW, etc.
- **notRepairedDamage** : if the car has a damage which is not repaired yet
- **dateCreated** : the date for which the ad at ebay was created
- **nrOfPictures** : number of pictures in the ad (unfortunately this field \* contains everywhere a 0 and is thus useless (bug in crawler) )
- **postalCode** : Area wise postal code
- **lastSeenOnline** : when the crawler saw this ad last online

```
df.describe() # Getting descriptive statistics
```

	price	yearOfRegistration	powerPS	kilometer	monthOfRegistration	nrOfPictures	postalCode
count	3.715280e+05	371528.000000	371528.000000	371528.000000	371528.000000	371528.0	371528.00000
mean	1.729514e+04	2004.577997	115.549477	125618.688228	5.734445	0.0	50820.66764
std	3.587954e+06	92.866598	192.139578	40112.337051	3.712412	0.0	25799.08247
min	0.000000e+00	1000.000000	0.000000	5000.000000	0.000000	0.0	1067.00000
25%	1.150000e+03	1999.000000	70.000000	125000.000000	3.000000	0.0	30459.00000
50%	2.950000e+03	2003.000000	105.000000	150000.000000	6.000000	0.0	49610.00000
75%	7.200000e+03	2008.000000	150.000000	150000.000000	9.000000	0.0	71546.00000
max	2.147484e+09	9999.000000	20000.000000	150000.000000	12.000000	0.0	99998.00000

# DATA EXPLORATION


## Examining distribution of all variables

- Analyzing the outliers using the box-plot
- Using histogram density by plotting a bell curve
- Removing the outliers using IQR and Manual removing technique

```
df_clean = df_clean[~((df_clean < (Q1-1.5 * IQR)) |(df_clean > (Q3 + 1.5 * IQR))).any(axis=1)]
```


```
df_clean['price'].plot.box()
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7faac4817940>



```
import numpy as np
import pandas as pd
from scipy.stats import norm
import matplotlib.pyplot as plt
df = pd.DataFrame({'price': np.random.normal(size = 100)})
df.price.plot(kind = 'hist', density = True)
range = np.arange(-3, 3, 0.001)
plt.plot(range, norm.pdf(range, 0, 1))
```

<matplotlib.lines.Line2D at 0x7f7e144e4208>





# DATA EXPLORATION


## Detecting missing values using plot

```
df_clean.isnull().sum()
```

```
price                0
vehicleType          8674
yearOfRegistration   0
gearbox              1866
powerPS              0
model                6556
kilometer            0
monthOfRegistration  0
fuelType             8419
brand                0
postalCode           0
dtype: int64
```

```
sns.heatmap(df_clean.isnull(), cbar=False)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f8b02cae400>



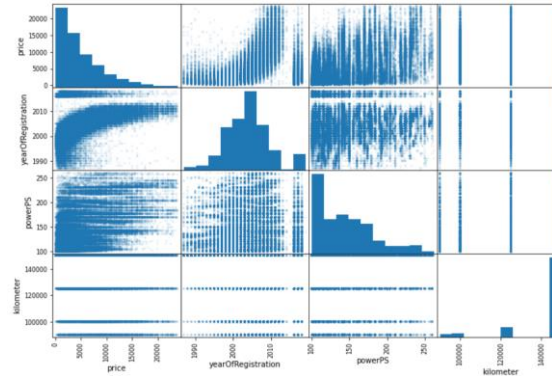

# DATA VISUALIZATION

Data was visualized using different plots

- Hexagonal bin plots
- Scatter matrix plot

```
df_clean.plot.hexbin(x = 'price', y = 'yearOfRegistration', gridsize = 25)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7df0b6d940>
```




# FEATURE SELECTION

- Chi Squared test

Features	Score
kilometer	8.093982e+08
postalCode	7.993496e+07
powerPS	1.354623e+06
model	3.623605e+05
brand	6.458665e+04
fuelType	9.892546e+03
monthOfRegistration	8.880040e+03
gearbox	6.577627e+03
vehicleType	5.282056e+03
yearOfRegistration	1.261282e+03

- Heat map

```
import seaborn as sns
# get correlations of each features in dataset
corrmat = df_clean.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(10,10))
g=sns.heatmap(df_clean[top_corr_features].corr(),annot=True,cmap="RdYlGn")
```



# Model Planning

- Variable selection

```
selectedFeatures =  
    [  
        'yearOfRegistration', 'powerPS', 'model', 'kilometer', 'monthOfRegistration', 'fuelType',  
        'brand', 'postalCode', 'vehicleType_0', 'vehicleType_1', 'vehicleType_2', 'vehicleType_3',  
        'vehicleType_4', 'vehicleType_5', 'vehicleType_6', 'vehicleType_7', 'gearbox_0', 'gearbox_1'  
    ]  
  
X = df[selectedFeatures]  
y = df['price']
```

- Model Selection

- Classification
- **Regression**
- Association Rules
- Text/Image/Video Analysis

# Model Building

- Splitting the dataset into three sets

- **Training set** - 80%
- **Validation set** - 10%
- **Testing set** - 10%

- Choosing the learning algorithm with validation set

```
from sklearn.model_selection import train_test_split

# Training set = 90%, Testing set = 10%, Validation set = 10%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=1)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.1, random_state=1)
```

```
start_time = time.time()

# Dictionary of pipelines and Regression types for ease of reference
pipe_dict = {0: 'Random Forest Regression', 1: 'Decision Tree Regressor', 2: 'Linear Regression', 3: 'Support Vector Regression'}
# Fit the pipelines
for pipe in pipelines:
    pipe.fit(X_train, y_train)

for i,model in enumerate(pipelines):
    pred = model.predict(X_val)
    print("{} Model Accuracy: {}".format(pipe_dict[i],r2_score(y_val, pred)* 100))

print("--- %s seconds ---" % (time.time() - start_time)) # Displaying the time in seconds

Random Forest Regression Model Accuracy: 84.81584796937892
Decision Tree Regressor Model Accuracy: 73.02462727685325
Linear Regression Model Accuracy: 56.095065680128066
Support Vector Regression Model Accuracy: 23.623095422425923
--- 379.20802187919617 seconds ---
```

# Model Building

## •Underfitting

- No underfitting

Regression Algorithms	Model Accuracy
Random Forest	97.85146144913742
Decision Tree	99.36287739304052

## •Overfitting

- Slightly overfitting

```
from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor().fit(X_train, y_train)
pred = rfr.predict(X_test)
print(r2_score(y_test, pred)* 100)
```

84.65840335933866

# Model Building

- **Hyper parameter tuning**

- Manually
- Grid search CV

Random forest reg	n_estimators	max_depth	max_features	min_samples_leaf	MAE	MSE	Accuracy	Time (s)
Manual tuning	270	14	18	11	1163.9861317221678	2964894.919932388	84.63549912181718	40.37581658363342
Grid Search CV	350	16	10	2	1108.0748354948025	2697919.915430242	86.01900100026474	39.78496479988098

# Model Building

## •L1 – Regularization

- **Lasso**

L1 - Regularization	Accuracy	Time taken to execute(s)
Lasso Regression - alpha = 10000	-0.004225239923694	0.01874542236328125
Lasso Regression - alpha = 1000	40.188960421043895	0.015836477279663086
Lasso Regression - alpha = 100	56.53602625199379	0.024311065673828125
Lasso Regression - alpha = 10	56.872121216214076	0.027382612228393555
Lasso Regression - alpha = 1	56.88514458882089	0.028135061264038086
Lasso Regression - alpha = 0.1	56.88610417645656	0.031087160110473633

## •L2 - Regularization

- **Ridge**

L2 - Regularization	Accuracy	Time taken to execute(s)
Ridge Regression - alpha = 10000	56.005916173458516	0.021957874298095703
Ridge Regression - alpha = 1000	56.862126442879735	0.024925947189331055
Ridge Regression - alpha = 100	56.88466608809992	0.017815828323364258
Ridge Regression - alpha = 10	56.88605350312452	0.02006673812866211
Ridge Regression - alpha = 1	56.88618326663119	0.018297672271728516
Ridge Regression - alpha = 0.1	56.88619615287895	0.021454572677612305



# Model Building

## •MSE train

```
from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor().fit(X_train, y_train)
pred = rfr.predict(X_train)

print("Mean Absolute Error is :", mean_absolute_error(y_train, pred))
print("-----")
print("Mean Squared Error is :", mean_squared_error(y_train, pred))
print("-----")
print("The R2 square value of Random Forest Regression is :",rfr.score(X_train, y_train)* 100)

Mean Absolute Error is : 426.01334188659473
-----
Mean Squared Error is : 410579.12151505775
-----
The R2 square value of Random Forest Regression is : 97.8539998222856
```

## •MSE test

```
from sklearn.ensemble import RandomForestRegressor
clf = RandomForestRegressor()
clf.fit(X_test, y_test)
pred = clf.predict(X_test)

print("Mean Absolute Error is :", mean_absolute_error(y_test, pred))
print("-----")
print("Mean Squared Error is :", mean_squared_error(y_test, pred))
print("-----")
print("The R2 square value of Random Forest Regression is :",clf.score(X_test, y_test)* 100)


Mean Absolute Error is : 453.6230156679866
-----
Mean Squared Error is : 457486.05584306625
-----
The R2 square value of Random Forest Regression is : 97.62924316153588
```

## •T-test

- $t = -124.323, p = 0.02$

# SOLUTION OVERVIEW

- Category of data falls under supervised machine learning: Regression Estimator



# OUTCOME

- Component Of Project

- Fully developed and Operational Cloud-based published website.  
(Copied from [UR Courses](#))

- Link:

[https://ucpps.herokuapp.com/ucpps\\_app/home.html](https://ucpps.herokuapp.com/ucpps_app/home.html)

The screenshot displays the 'Used Car Price Prediction System' (UCPPS) web application. The interface is titled 'UCPPS' in a blue header. Below the header, a subtitle reads 'How much is your car worth?, Enter the specifications of the car and get the price'. The form contains several input fields and dropdown menus: 'Name of the car' (text input), 'Brand of the car' (dropdown menu showing 'alfa\_romeo'), 'Model of the car' (dropdown menu showing '3 er'), 'Total Kilometers' (text input), 'Type of the car' (dropdown menu showing 'Limousine'), 'Gearbox of the car' (dropdown menu showing 'Manual'), 'Power of the car (PS)' (text input), 'Fueltype of the car' (dropdown menu showing 'Petrol'), 'Year of Registration' (dropdown menu showing '1995'), 'Month of Registration' (dropdown menu showing '1'), and 'Postal Code' (text input). A 'Submit' button is located at the bottom of the form. Callouts provide additional information: 'Text field to enter the input' points to the 'Name of the car' field; 'Drop down menu to choose an input' points to the 'Fueltype of the car' dropdown; 'On clicking this button, the price will be generated' points to the 'Submit' button; and 'Feedback option' points to a link at the bottom of the page that says 'Not satisfied with the results? Click here to give your feedback'. The footer of the application shows the title 'Used Car Price Prediction System', the date 'Mon Apr 6 12:00:21 2020', and the feedback link.

# TOOLS

- [Google Colab](#)
- [Anaconda Jupyter Notebook](#)
- [Python Programming \(Obviously\)](#)
- Python Libraries ([Pandas](#), [NumPy](#), [Matplotlib](#), [scikit learn](#))
- [Django](#) Web framework ([HTML](#), [CSS](#), [Bootstrap](#))
- [GitHub](#)
- [Visual Studio Code](#)
- [Heroku](#)


# TEAM ROLES

- Data Collection, data understanding
  - Model Design, model evaluation
  - Code Documentation
  - Deployment and building a functional website.
- 
- 2 new things that took place recently
    - Uploaded my first YouTube video (Live coding) about code deployment
    - Deployed my first project on to the cloud



Image Credits: [Medium](#)

# TIMELINE



# REFERENCES

[1] Manashty, D. (2020). *Data Science Fundamentals - Chapter 1*. Presentation, University of Regina, Canada.

[2] *The all-new BMW 5-series (G30) launched – All You Need to Know*. (2020). [Image]. Retrieved 8 April 2020, from <http://kensomuse.com/blog/2017/03/30/new-bmw-5-series-g30-launched-need-know/>

[3] Leka, O. Used Cars Data - dataset by data-society. Retrieved 8 April 2020, from <https://data.world/data-society/used-cars-data>

Thank you